

Lecture #4

➤ HeapSort

Prof. Pietro S. Oliveto

Department of Computer Science and Engineering
Southern University of Science and Technology (SUSTech)

olivetop@sustech.edu.cn
<https://faculty.sustech.edu.cn/olivetop>

➤ Aims of this lecture

- To introduce the **HeapSort** algorithm.
- To show how a **clever data structure**, a **heap**, can lead to a **fast** and **in place** sorting algorithm
 - **In place: $O(1)$ additional space.**
- To **practice the design and analysis of algorithms.**

Reading: Chapter 6

➤ Idea behind HeapSort

- Idea:
 - Find the largest element.
 - Move it to the end of the array (put another one in its place).
 - Repeat with remaining elements.
- Like SelectionSort but ...
 - SelectionSort compares lots of elements to find the largest.
 - **Can we store knowledge gained from these comparisons for the future?**
 - Use this knowledge to make future iterations faster!

➤ Use your imagination...

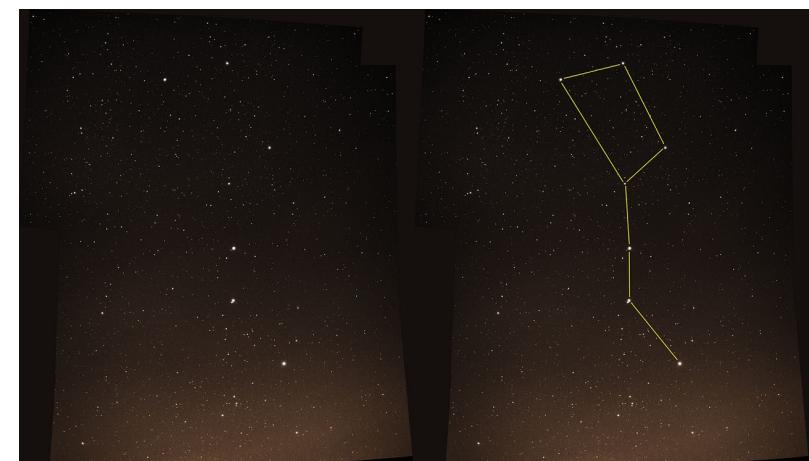
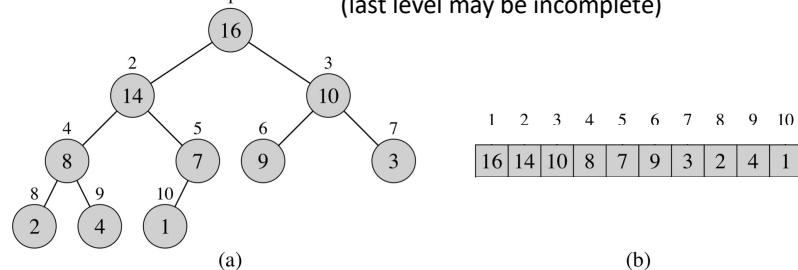


Photo : Thomas Bresson

➤ A Heap

- Essentially an array **imagined** as being a **binary tree**!
 - Elements are arranged row by row from left to right.


(last level may be incomplete)



- Navigate through the array/imaginary tree using these operations:

- $\text{Parent}(i) = \left\lfloor \frac{i}{2} \right\rfloor$ (“floor of $i/2$ ”), $\text{Left}(i) = 2i$, $\text{Right}(i) = 2i + 1$

数组的索引顺序就是节点的“排队”顺序。一个节点的子节点，只有在

CSE217: Data Structures & Algorithm Analysis

5

第*i*个节点的子节点：前面有*i*-1个节点作为父节点，生成 $2i-2$ 个子节点，加上最前面的第一个根节点，一共 $2i-1$ 个节点，后面就是第*i+1*和

► Procedures (what do we need)

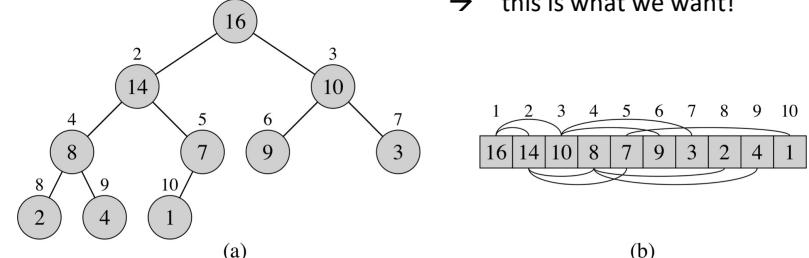
1. **Build-Max-Heap**: produces a Max-Heap from an unordered array
 2. **Max-Heapify**: maintains the max-heap property once the maximum has been removed
 3. **HeapSort**: sorts an array in place

- New variable $A.\text{heap-size}$ indicates how many elements of A are stored in a heap: $0 \leq A.\text{heap-size} \leq A.\text{length}$.
 - Decreasing $A.\text{heap-size}$ by 1 effectively removes the last element from the heap (we imagine a heap without it)
 - There are analogous operations for min-heaps:
Min-Heapify and Build-Min-Heap.



➤ Heap Properties

- **Max-heap property:** for every node other than the root, the parent is no smaller than the node, $A[\text{Parent}(i)] \geq A[i]$.
 - In a max-heap, the **root** always stores a **largest** element.
→ this is what we want!



- **Min-heap property:** for every node other than the root, the parent is no larger than the node, $A[\text{Parent}(i)] \leq A[i]$.

CSE217: Data Structures & Algorithm Analysis

6

➤ Procedures (what do we need)

1. **Build-Max-Heap**: produces a Max-Heap from an unordered array
 2. **Max-Heapify**: maintains the max-heap property once the maximum has been removed
 3. **HeapSort**: sorts an array in place

- New variable $A.\text{heap-size}$ indicates how many elements of A are stored in a heap: $0 \leq A.\text{heap-size} \leq A.\text{length}$.
 - Decreasing $A.\text{heap-size}$ by 1 effectively removes the last element from the heap (we imagine a heap without it)
 - There are analogous operations for min-heaps:
Min-Heapify and Build-Min-Heap.

CSE217: Data Structures & Algorithm Analysis

7

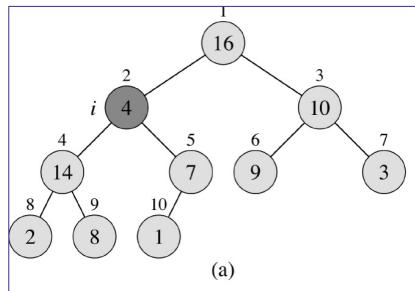
CSE217: Data Structures & Algorithm Analysis

8

注意这边，Max-Heapify是一个用来递归构造一个Max-Heap的方法，也就是from bottom to up建立。上面的1、2、3步并不是步骤，2 Max-Heapify也被包含在1中的build过程中！

➤ Max-Heapify(A, i)

- Assumes subtrees Left(i) and Right(i) are max-heaps**, but max-heap property might be violated in root of subtree at i .
 - “Subtree x”: the part of the tree including x and everything below.
- Lets the value at $A[i]$ “float down” if necessary, to restore max-heap property at i
- At the end of Max-Heapify the subtree at i is a max-heap.



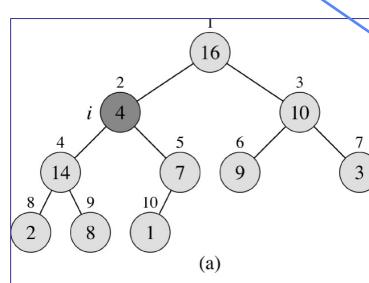
CSE217: Data Structures & Algorithm Analysis

9

- i 是一个叶子节点： $i=2i$ 会超过堆的边界，也就是没有子节点。
- 在堆排序过程中，堆的大小动态缩小时，比如 $i = \text{LEFT}(5) = 10 \leq 9$ ，会显示 False，这时候子节点是上一回合中最大的那个数，已不在当前回合的范围。

➤ Max-Heapify: informal and in pseudocode

- Compare $A[i]$ with all existing children
- If largest child is larger than $A[i]$, swap and recurse on child



CSE217: Data Structures & Algorithm Analysis

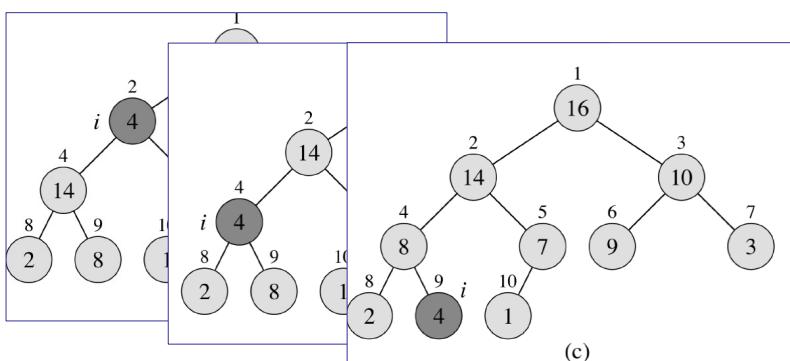
10

```
MAX-HEAPIFY( $A, i$ )
1:  $l = \text{Left}(i)$ 
2:  $r = \text{Right}(i)$ 
3: if  $l \leq A.\text{heap-size}$  and  $A[l] > A[i]$  then
4:   largest =  $l$ 
5: else
6:   largest =  $i$ 
7: if  $r \leq A.\text{heap-size}$  and  $A[r] > A[\text{largest}]$  then
8:   largest =  $r$ 
9: if largest  $\neq i$  then
10:  exchange  $A[i]$  with  $A[\text{largest}]$ 
11:  MAX-HEAPIFY( $A, \text{largest}$ )
```

现在的largest这个位置还是前面的那个位置，现在这个位置上的数字变成了之前i上的数字。我们对这个子树再重复验证它的root是不是还有上面的violation。

➤ Max-Heapify: Example

- Compare $A[i]$ with all existing children
- If largest child is larger than $A[i]$, swap and recurse on child



CSE217: Data Structures & Algorithm Analysis

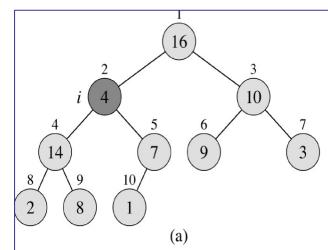
11

注意: h实际上与n有关的数, 所以并不能 $O(h)=O(1)$ 。

➤ Runtime of Max-Heapify

- Define the height of a node as the longest number of simple downward edges from the node to a leaf.
- Leaf: a node without children.
- Max-Heapify takes constant time, $\Theta(1)$, on each level.
- Running time of Max-Heapify on a node of height h is $O(h)$.
- It's not $\Omega(h)$ as Max-Heapify may stop early, e.g. if heap-property holds at i .
- For leaves $h = 0$ and the time is $O(1)$.

```
MAX-HEAPIFY( $A, i$ )
1:  $l = \text{Left}(i)$ 
2:  $r = \text{Right}(i)$ 
3: if  $l \leq A.\text{heap-size}$  and  $A[l] > A[i]$  then
4:   largest =  $l$ 
5: else
6:   largest =  $i$ 
7: if  $r \leq A.\text{heap-size}$  and  $A[r] > A[\text{largest}]$  then
8:   largest =  $r$ 
9: if largest  $\neq i$  then
10:  exchange  $A[i]$  with  $A[\text{largest}]$ 
11:  MAX-HEAPIFY( $A, \text{largest}$ )
```



CSE217: Data Structures & Algorithm Analysis

12

➤ Bounding the height of a heap

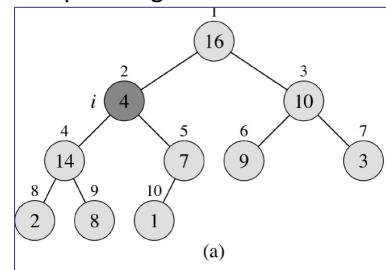
- **Claim:** the height of a heap = height of the root is at most $\log n$.

- **Proof:** the number n of elements in a heap of height h is

- Doubling on each level
 - At least 1 node on the last level
 - Hence in total at least

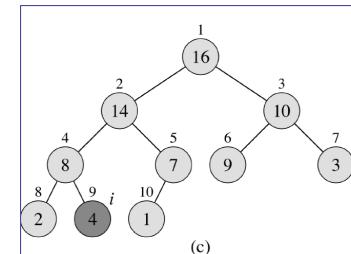
$$1 + 2 + 4 + \cdots + 2^{h-1} + 1 = 2^h$$

(we used $\sum_{i=0}^{k-1} 2^i = 2^k - 1$)



- So size and height are related as $n \geq 2^h \Leftrightarrow \log n \geq h$
 - “the height of the root is at most $\log n$ ”
 - So the runtime of Max-Heapify is $O(\log n)$

► Max-Heapify: Correctness



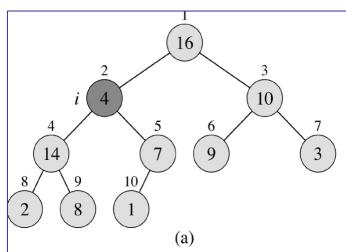
```

MAX-HEAPIFY( $A, i$ )
1:  $l = \text{Left}(i)$ 
2:  $r = \text{Right}(i)$ 
3: if  $l \leq A.\text{heap-size}$  and  $A[l] > A[i]$  then
4:   largest =  $l$ 
5: else
6:   largest =  $i$ 
7: if  $r \leq A.\text{heap-size}$  and  $A[r] > A[\text{largest}]$  then
8:   largest =  $r$ 
9: if largest  $\neq i$  then
10:   exchange  $A[i]$  with  $A[\text{largest}]$ 
11:   MAX-HEAPIFY( $A, \text{largest}$ )

```

- By induction (on the height):
 - Base case: height = 0 (i is a leaf)
 - Then left(i) and right(i) are larger than A.heap-size and the algorithm returns a heap!

➤ Max-Heapify: Correctness



- By induction (on the height):
10: exchange $A[i]$ with $A[\text{largest}]$
11: MAX-HEAPIFY(A , largest)
 - **Inductive case:** assume it works for height $h = i - 1$ and show it works for $h = i$
 - Then the algorithm swaps $A[i]$ with the larger between $\text{Left}(i)$ and $\text{Right}(i)$ (if any) and one subtree was already a heap and the other will be by inductive hypothesis.

➤ Procedures (what do we need)

1. **Build-Max-Heap**: produces a Max-Heap from an unordered array

- 2. Max-Heapify:** maintains the max-heap property once the maximum has been removed ✓

- 3. HeapSort:** sorts an array in place

➤ Building a Heap

- Idea: use Max-Heapify repeatedly to create a heap.
- Which order of nodes: top-down or bottom-up?
- Answer: **bottom-up** – Max-Heapify assumes Left(i) and Right(i) are heaps. Top-down wouldn't work, bottom-up does.
- Note: nodes in $A \left[\left(\left\lfloor \frac{n}{2} \right\rfloor + 1 \right), \dots, n \right]$ are all leaves. Leaves are max-heaps, so no work required.

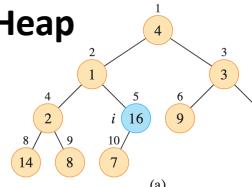
BUILD-MAX-HEAP(A, n)

```

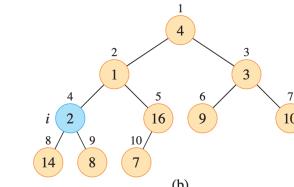
1  A.heap-size = n
2  for i = ⌊n/2⌋ downto 1
3      MAX-HEAPIFY(A, i)
```

$A [4 | 1 | 3 | 2 | 16 | 9 | 10 | 14 | 8 | 7]$

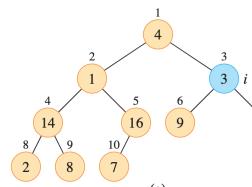
➤ Build-Max-Heap



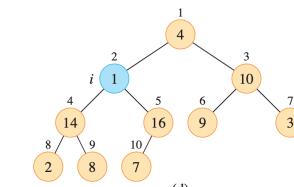
(a)



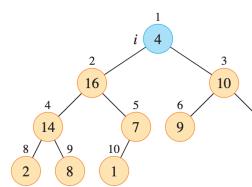
(b)



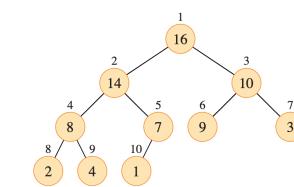
(c)



(d)



(e)



(f)

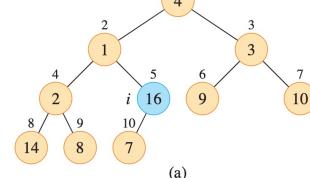
➤ Correctness of Build-Max-Heap

BUILD-MAX-HEAP(A, n)

```

1  A.heap-size = n
2  for i = ⌊n/2⌋ downto 1
3      MAX-HEAPIFY(A, i)
```

- Loop invariant:** At the start of each iteration i of the for loop, each node $i + 1, i + 2, \dots, n$ is the root of a max-heap.
- Initialisation:** true for leaves $\left\lfloor \frac{n}{2} \right\rfloor + 1, \dots, n$.



(a)

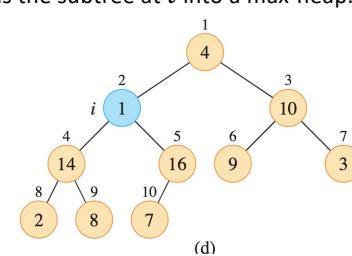
➤ Correctness of Build-Max-Heap

BUILD-MAX-HEAP(A, n)

```

1  A.heap-size = n
2  for i = ⌊n/2⌋ downto 1
3      MAX-HEAPIFY(A, i)
```

- Loop invariant:** At the start of each iteration i of the for loop, each node $i + 1, i + 2, \dots, n$ is the root of a max-heap.
- Maintenance:** by loop invariant, all children of i are roots of max-heaps (as their numbers are larger than i).
Then Max-Heapify(A, i) turns the subtree at i into a max-heap.



(d)

➤ Correctness of Build-Max-Heap

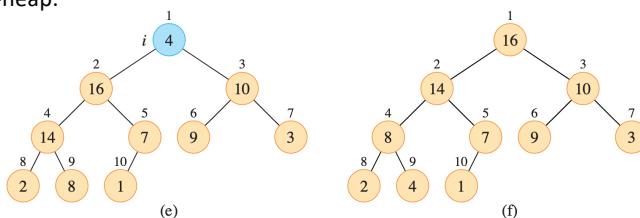
BUILD-MAX-HEAP(A, n)

```

1    $A.heap\_size = n$ 
2   for  $i = \lfloor n/2 \rfloor$  downto 1
3       MAX-HEAPIFY( $A, i$ )

```

- Loop invariant:** At the start of each iteration i of the for loop, each node $i + 1, i + 2, \dots, n$ is the root of a max-heap.
- Termination:** the loop terminates at $i = 0$, hence node 1 is the root of a max-heap.



➤ Runtime of Build-Max-Heap

- The **height of a heap** = height of the root is at most $\log n$.
- So all nodes have height at most $\log n$.
- Every call to Max-Heapify takes time $O(\log n)$.
- Build-Max-Heap calls Max-Heapify $O(n)$ times.
- Total time is at most $O(n) \cdot O(\log n) = O(n \log n)$.
 - The time can be improved to $O(n)$ since most nodes have small height.
 - $O(n \log n)$ is sufficient for us, though.

➤ Refined Analysis of Build-Max-Heap

- Observation: most nodes have small height**

One can show: there are at most $\left\lceil \frac{n}{2^{h+1}} \right\rceil$ nodes of height h .

$O(\log n)$ time bound is correct, but crude for most nodes.

A better bound:

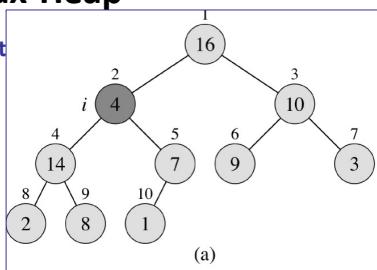
$$\sum_{h=1}^{\lfloor \log n \rfloor} \left\lceil \frac{n}{2^{h+1}} \right\rceil O(h) = O\left(n \sum_{h=1}^{\lfloor \log n \rfloor} \frac{h}{2^h}\right) = O\left(n \sum_{h=1}^{\infty} \frac{h}{2^h}\right) = O(n)$$

as the infinite series of $\frac{h}{2^h}$ is 2.

- 1st equality, we used that: $[x] \leq 2x$ for $x \geq 1/2$

\Rightarrow for $h \leq \log n$, $\frac{n}{2^{h+1}} \geq 1/2$ because $n \geq 2^h$ (see slide 13)

- 2nd equality, we used that $\sum_{k=0}^{\infty} kx^k = \frac{x}{(1-x)^2}$ for $|x| < 1$



(a)

➤ Procedures (what do we need)

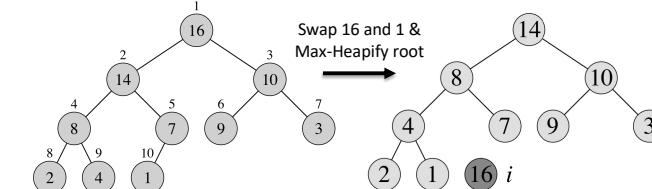
1. Build-Max-Heap: produces a Max-Heap from an unordered array ✓

2. Max-Heapify: maintains the max-heap property once the maximum has been removed ✓

3. HeapSort: sorts an array in place

➤ HeapSort

- Ideas:



- Build a max-heap, such that the root contains largest element.
- Swap the root with the last element of the heap/array.
- Discard the last element from the heap by reducing heap.size.
(We simply imagine a smaller heap.)
- Call Max-Heapify($A, 1$) to restore heap property at the root.

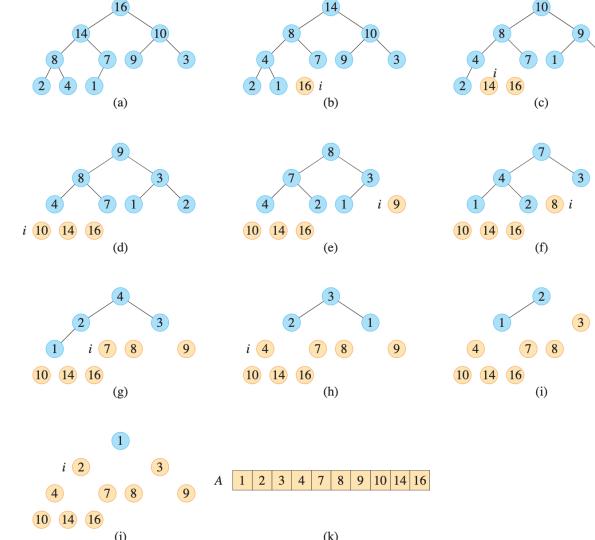
HEAPSORT(A)

```

1: BUILD-MAX-HEAP( $A$ )
2: for  $i = A.length$  downto 2 do
3:   exchange  $A[1]$  with  $A[i]$ 
4:    $A.heap-size = A.heap-size - 1$ 
5:   MAX-HEAPIFY( $A, 1$ )

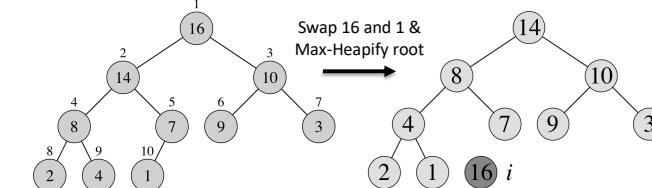
```

➤ HeapSort: Example



➤ HeapSort

- Ideas:



- Build a max-heap, such that the root contains largest element.
- Swap the root with the last element of the heap/array.
- Discard the last element from the heap by reducing heap.size.
(We simply imagine a smaller heap.)
- Call Max-Heapify($A, 1$) to restore heap property at the root.

HEAPSORT(A)

```

1: BUILD-MAX-HEAP( $A$ )
2: for  $i = A.length$  downto 2 do
3:   exchange  $A[1]$  with  $A[i]$ 
4:    $A.heap-size = A.heap-size - 1$ 
5:   MAX-HEAPIFY( $A, 1$ )

```

Runtime:

$$\begin{aligned}
 & O(n \log n) \\
 & + (n - 1) \cdot O(\log n) \\
 & = O(n \log n)
 \end{aligned}$$

➤ Correctness of HeapSort

Loop Invariant: “At the start of each iteration of the for loop of lines 2-5, the subarray $A[1..i]$ is a max-heap containing the i smallest elements of $A[1..n]$, and the subarray $A[i+1..n]$ contains the $n-i$ largest elements of $A[1..n]$, sorted.”

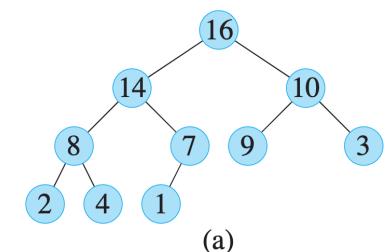
- Initialization:** The subarray $A[i+1..n]$ is empty, thus the invariant holds.

HEAPSORT(A)

```

1: BUILD-MAX-HEAP( $A$ )
2: for  $i = A.length$  downto 2 do
3:   exchange  $A[1]$  with  $A[i]$ 
4:    $A.heap-size = A.heap-size - 1$ 
5:   MAX-HEAPIFY( $A, 1$ )

```

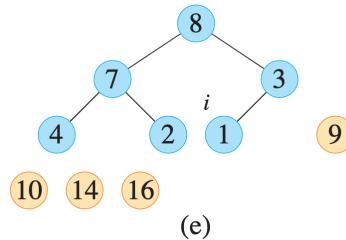


➤ Correctness of HeapSort

Loop Invariant: “At the start of each iteration of the for loop of lines 2-5, the subarray $A[1..i]$ is a max-heap containing the i smallest elements of $A[1..n]$, and the subarray $A[i+1..n]$ contains the $n-i$ largest elements of $A[1..n]$, sorted.”

Maintenance: $A[1]$ is the largest element in $A[1..i]$ and it is smaller than the elements in $A[i+1..n]$. When we put it in the i th position, then $A[i..n]$ contains the largest elements, sorted. Decreasing the heap size and calling Max-Heapify turns $A[1..i-1]$ into a max-heap. Decrementing i sets up the invariant for the next iteration.

```
HEAPSORT( $A$ )
1: BUILD-MAX-HEAP( $A$ )
2: for  $i = A.length$  downto 2 do
3:   exchange  $A[1]$  with  $A[i]$ 
4:    $A.heap-size = A.heap-size - 1$ 
5:   MAX-HEAPIFY( $A, 1$ )
```



CSE217: Data Structures & Algorithm Analysis

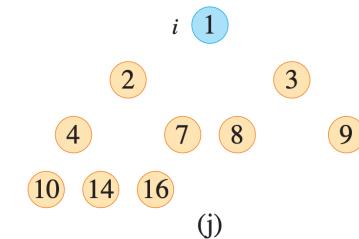
29

➤ Correctness of HeapSort

Loop Invariant: “At the start of each iteration of the for loop of lines 2-5, the subarray $A[1..i]$ is a max-heap containing the i smallest elements of $A[1..n]$, and the subarray $A[i+1..n]$ contains the $n-i$ largest elements of $A[1..n]$, sorted.”

- **Termination:** After the loop $i=1$. This means that $A[2..n]$ is sorted and $A[1]$ is the smallest element in the array, which makes the array sorted.

```
HEAPSORT( $A$ )
1: BUILD-MAX-HEAP( $A$ )
2: for  $i = A.length$  downto 2 do
3:   exchange  $A[1]$  with  $A[i]$ 
4:    $A.heap-size = A.heap-size - 1$ 
5:   MAX-HEAPIFY( $A, 1$ )
```



CSE217: Data Structures & Algorithm Analysis

30

➤ Summary

- HeapSort sorts in place in time $O(n \log n)$.
 - Building a Heap in time $O(n)$.
 - Extracting the largest element and restoring the heap-property in total time $O(n \log n)$.
- The use of appropriate **data structures** can speed up computation (in contrast to SelectionSort).
 - The heap “memorises” information about comparisons of elements.
 - The heap is imaginary, no objects/pointers required!
- Heaps also play a role in **Priority Queues**.

CSE217: Data Structures & Algorithm Analysis

31

关键区别：重建 (Rebuild) vs. 修复 (Repair)

让我们来对比一下这两种思路：

1. 浪费时间的思路 (重建整个堆)

我们有一个 n 个元素的最大堆。取出最大值 $A[1]$ 。

现在我们剩下 $n-1$ 个元素，它们是 $A[2], A[3], \dots, A[n]$ 。

(浪费时间的步骤) 把这 $n-1$ 个元素当成一个全新的、无序的数组，然后调用 Build-Max-Heap 算法，花费 $O(n)$ 的时间重新把它们组织成一个堆。

重复这个过程...

如果按照这个思路，每次取出一个元素都要花费 $O(n)$ 的时间来重建堆，总时间复杂度会是 $O(n^2)$ ，这和低效的选择排序、冒泡排序就没什么区别了。

2. 高效的思路 (Heapsort的实际做法：仅修复)

Heapsort的实际做法是：在取出最大值后，堆的结构几乎是完好的，只有一个地方出了问题。我们只需要修复那个小问题就行了。

发生了什么？交换：我们将堆顶 $A[1]$ （最大值）和堆的最后一个元素 $A[\text{heap.size}]$ 交换。缩小： $\text{heap.size}--$ 。

现在的局面是什么？根节点 (Root): $A[1]$ 现在存放的是原来那个很小的叶子节点的值。这导致最大堆的性质在根部被破坏了。 $A[1]$ 很可能比它的子节点 $A[2]$ 和 $A[3]$ 要小。其他部分 (The Rest of the Tree): 根节点的两个子树（以 $A[2]$ 和 $A[3]$ 为根的子树）本身仍然是完美的最大堆！因为我们根本没有动它们内部的结构。

我们要做的工作：我们不需要重建一切。我们只需要把那个被安插在CEO位置上的“实习生”（原来那个小的值）一路“下沉”或“筛选”到它合适的位置就行了。这个操作就是 Max-Heapify。

Max-Heapify($A, 1$) 的工作效率：

它从根节点 $i=1$ 开始。

它将 $A[1]$ 的值和它的子节点比较，然后和较大的那个子节点交换。

现在，那个小的值下沉了一层。它可能又破坏了下一层的堆性质。

所以 Max-Heapify 会递归地沿着一条路径继续往下走，直到这个元素不再小于它的子节点，或者它自己变成了叶子节点。

这个过程有多快？

这个“下沉”的路径最长能有多长？就是这个堆的高度。

一个包含 k 个元素的完全二叉树，它的高度大约是 $\log_2(k)$ 。

所以，每一次“修复”操作的时间复杂度仅仅是 $O(\log n)$ 。